

# Event Driven Architecture w/ Apache Kafka and Spring Cloud Stream

...

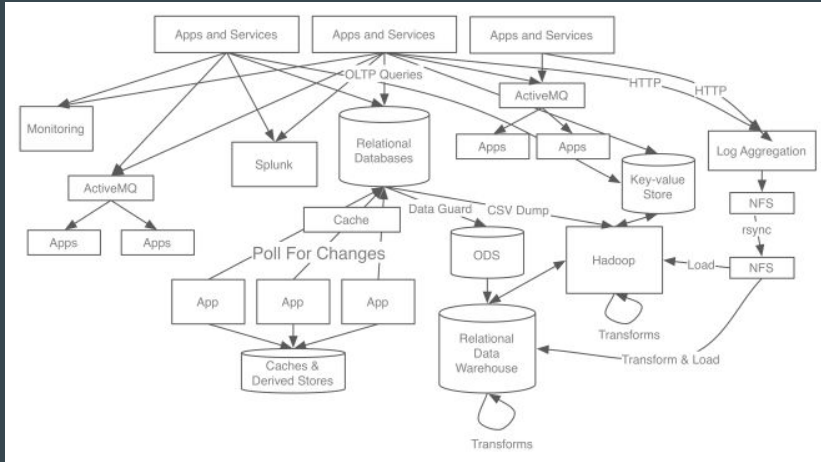
Indy Java User Group Feb 2019

# About Me

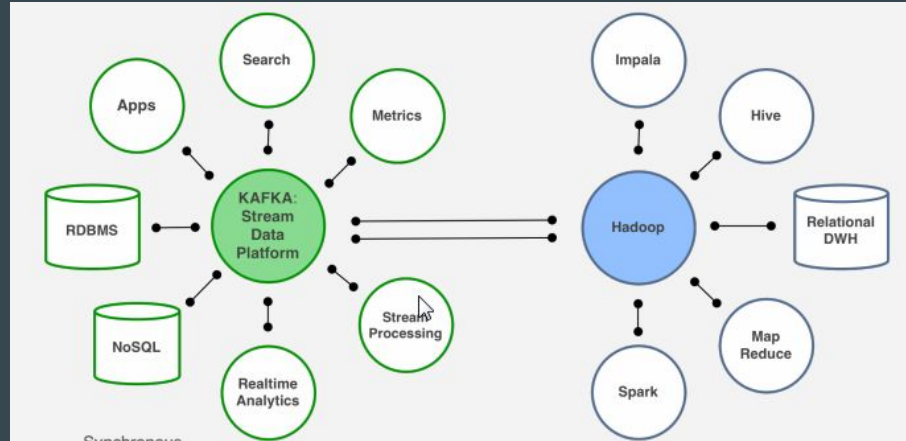
Dan Gradl  
Principal Engineer @ Finish Line (JD Sports)  
e: [dangradl@gmail.com](mailto:dangradl@gmail.com)



# Streaming Data Architecture



TO



# Key Patterns

**Source** - Stream data in from some location (file, database, etc.)

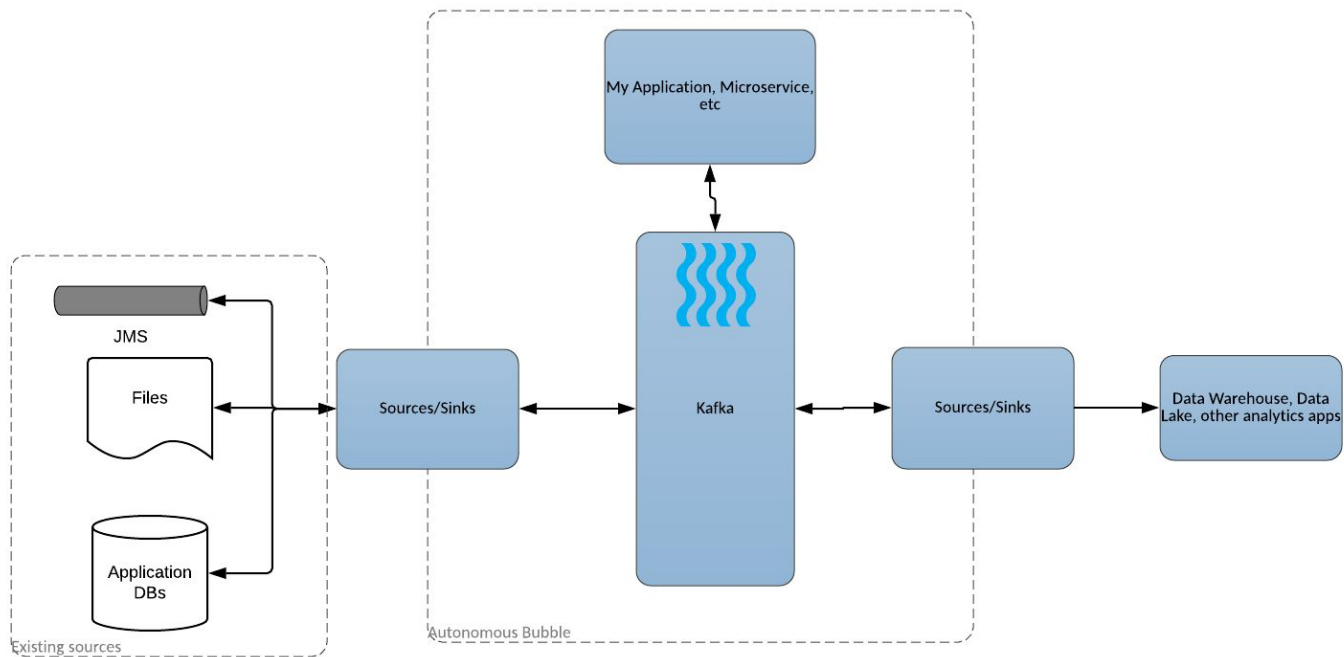
**Sink** - Stream data out to some location (file, database, etc.)

**Processor** - Process stream data from one stream and output to another including:

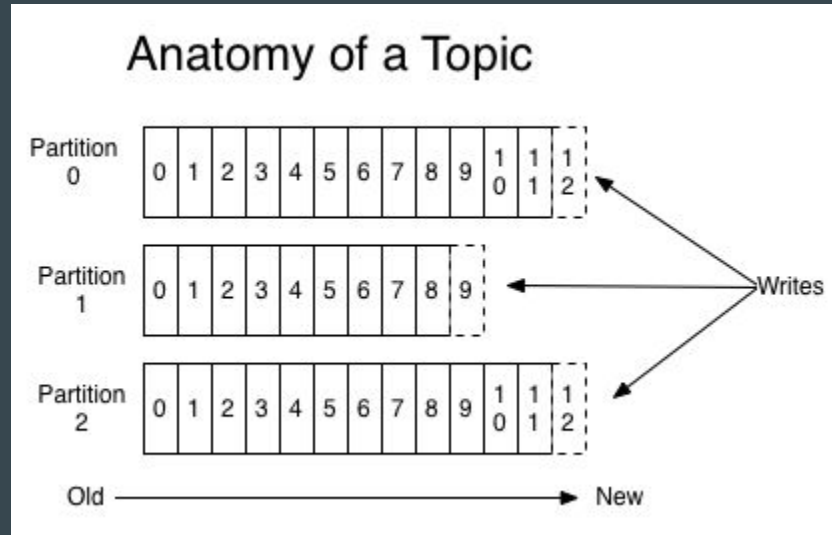
- Aggregation, filtering, and transforming of data

- Joining two streams and outputting to another

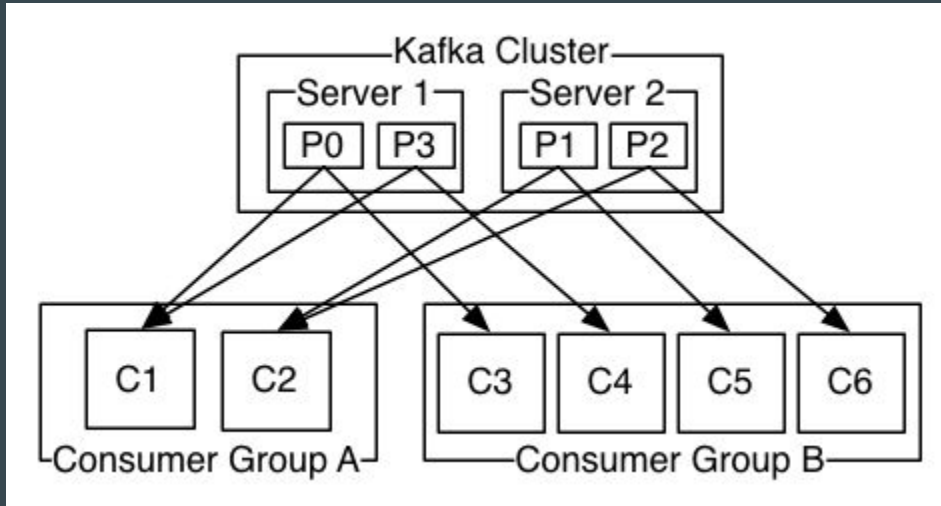
# Autonomous Bubble



# Kafka Topics



# Kafka Consumers



# Kafka features that support streaming architecture

- Low latency, high throughput
- Scalability - in particular partitioning enables parallelization of the stream processing for speed
- Flexible topic data retention
  - Size
  - Time
  - Compaction
- Support for different data payloads: e.g. Avro, JSON, Plain Text
- Replication for resilience and high availability
- Consumer controlled offset
- Kafka Connect and Kafka Streams



# Concrete Pattern Implementations

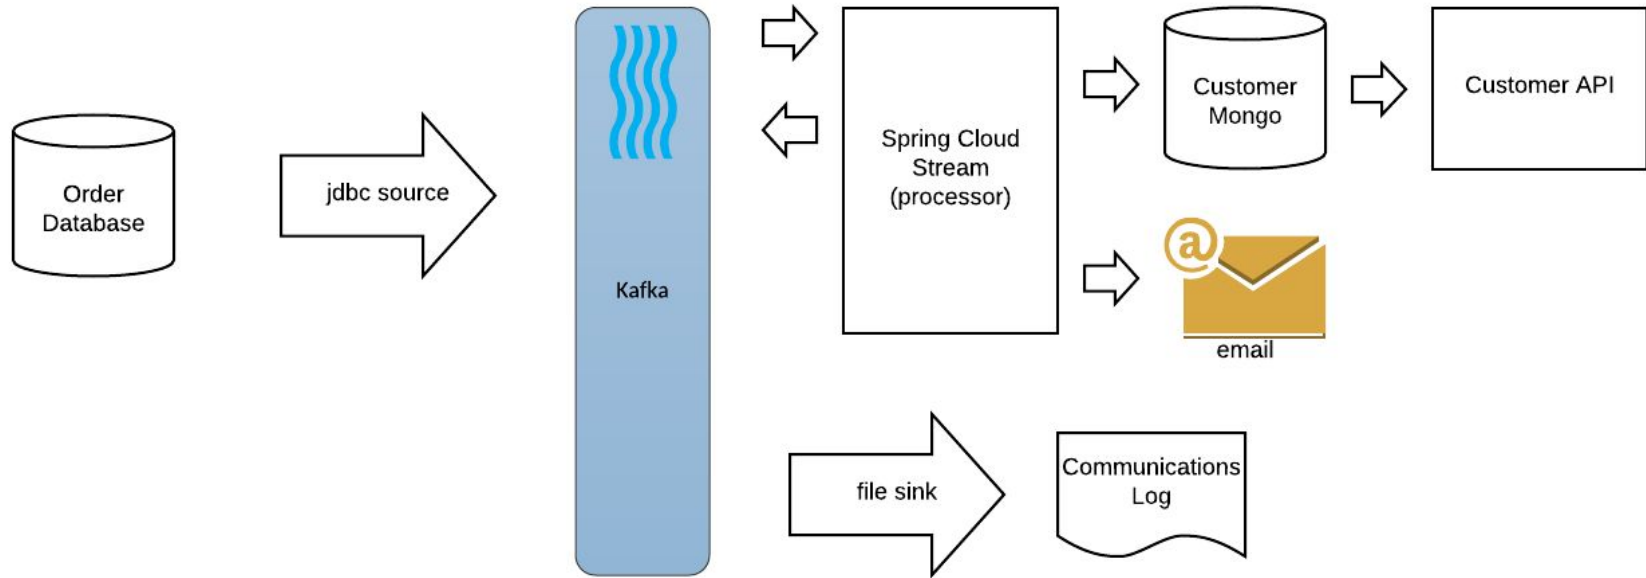
## Sources and Sinks

- Kafka Connect and Confluent connectors
- Spring Cloud Stream and associated App Starters
- Open source and commercial tools such as Striim and Streamsets
- Kafka console commands (poor man's)

## Processors

- Kafka Streams
- Spring Cloud Stream

# Patterns and Demo Use Case



# Order Generator and JDBC Source

Created a simple project that simulates a website or other application creating new orders in a postgres database (every 30 seconds).

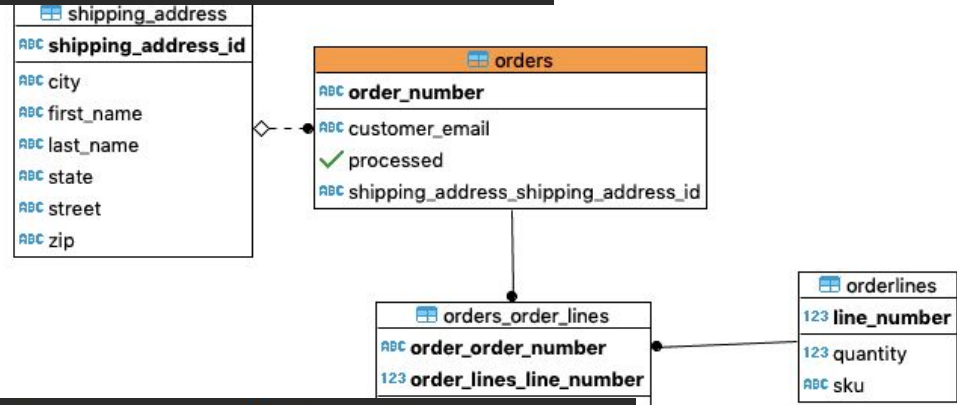
## Spring Boot App Starters

- Prepackaged sources, sinks and processors
- Configurable
- Simple to use

Source	Processor	Sink
file	aggregator	aggregate-counter
ftp	bridge	cassandra
gemfire	filter	counter
gemfire-cq	groovy-filter	field-value-counter
http	groovy-transform	file
jdbc	header-enricher	ftp
jms	httpClient	gemfire
load-generator	pmml	gpfdist
loggregator	python-http	hdfs
mail	python-ivythor	hdfs-dataset

# Orders - Polling

```
select o.order_number,  
       o.customer_email,  
       sa.first_name,  
       sa.last_name,  
       sa.street,  
       sa.city,  
       sa.state,  
       sa.zip  
from orders o  
      inner join shipping_address sa on o.shipping_address_shipping_address_id = sa.shipping_address_id  
where o.processed = false
```



```
update orders set processed=true where order_number in (:order_number)
```

# Running the App Starter

```
java -jar jdbc-source-kafka-2.1.0.RC1.jar \  
--server.port=0 \  
--spring.cloud.stream.bindings.output.destination=order_customers \  
--jdbc.query="select o.order_number,o.customer_email,sa.first_name,last_name,sa.street,sa.city,sa.state,sa.zip from ord \  
--jdbc.update="update orders set processed=true where order_number in (:order_number)" \  
--spring.datasource.url=jdbc:postgresql://localhost/postgres \  
--spring.datasource.username=postgres \  
--spring.datasource.password=mypassword
```

Database connection info

Query to retrieve data from DB and another to mark what has been polled

Destination binding for Kafka

# JDBC Source Output

```
{  
  "order_number": "368",  
  "customer_email": "idell.bednar@example.com",  
  "first_name": "Genie",  
  "last_name": "Rice",  
  "street": "2585 Renner Walks",  
  "city": "Cristville",  
  "state": "MD",  
  "zip": "10154-7785"  
}
```

# Initialize Your Project

<https://start.spring.io/>

## SPRING INITIALIZR

bootstrap your application now

Generate a Gradle Project with Java and Spring Boot 2.1.3

### Project Metadata

Artifact coordinates

Group

Artifact

### Dependencies

Add Spring Boot Starters and dependencies to your application

Search for dependencies

Selected Dependencies

Cloud Stream MongoDB Web Mail

Generate Project

# Additional Setup

Add the kafka binder to the build.gradle

```
dependencies {  
    implementation 'org.springframework.boot:spring-boot-starter-actuator'  
    implementation 'org.springframework.cloud:spring-cloud-stream'  
    implementation 'org.springframework.cloud:spring-cloud-stream-binder-kafka'  
    implementation 'org.springframework.data:spring-data-mongodb'  
    implementation 'org.springframework.boot:spring-boot-starter-web'  
    implementation 'org.springframework.boot:spring-boot-starter-mail'  
    implementation 'de.flapdoodle.embed:de.flapdoodle.embed.mongo'  
    testImplementation 'org.springframework.boot:spring-boot-starter-test'  
    testImplementation 'org.springframework.cloud:spring-cloud-stream-test-support'  
}
```

Point application.yml to our  
local (docker) kafka

```
spring:  
  kafka:  
    bootstrap-servers: localhost:9092
```



# Customer Repository and API

```
@Document(collection = "customers")
public class Customer {
    private String firstName;
    private String lastName;
    private String email;

    public Customer(String firstName, String lastName, String email) {...}

    public String getFirstName() { return firstName; }

    public Customer setFirstName(String firstName) {...}

    public String getLastName() { return lastName; }

    public Customer setLastName(String lastName) {...}

    public String getEmail() { return email; }

    public Customer setEmail(String email) {...}
}
```

```
@RestController
@RequestMapping(produces = MediaType.APPLICATION_JSON_UTF8_VALUE)
public class CustomerController {
    @Autowired
    CustomerRepository customerRepository;

    @GetMapping(path = "/customers")
    public Iterable<Customer> getCustomers() { return customerRepository.findAll(); }
}
```

```
public interface CustomerRepository extends CrudRepository<Customer, String> {
}
```

# Our First Stream - Sink

```
@EnableBinding(Sink.class)
public class OrderToCustomerProcessor {

    @Autowired
    CustomerRepository customerRepository;

    @StreamListener(Sink.INPUT)
    public void createCustomerFromOrder(OrderCustomerMessage orderCustomerMessage){
        Customer customer = new Customer(orderCustomerMessage.getFirstName(),
                                         orderCustomerMessage.getLastName(),
                                         orderCustomerMessage.getCustomerEmail());

        customerRepository.save(customer);
    }
}
```

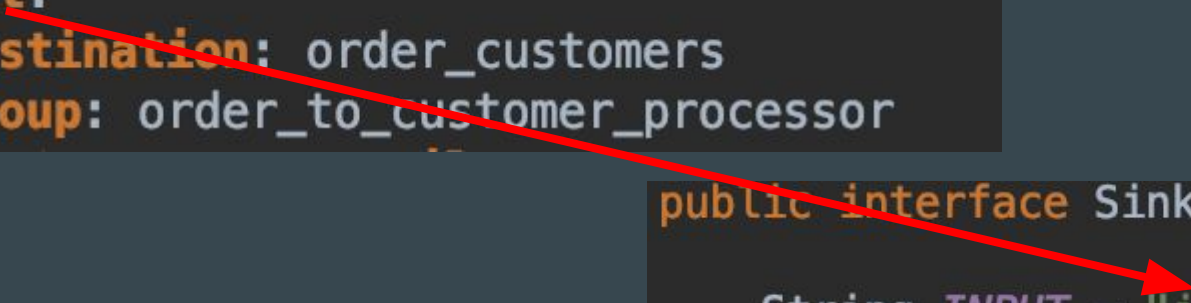
# OrderCustomerMessage

```
public class OrderCustomerMessage {  
  
    @JsonAlias("order_number")  
    private String orderNumber;  
  
    @JsonAlias("customer_email")  
    private String customerEmail;  
  
    @JsonAlias("first_name")  
    private String firstName;  
  
    @JsonAlias("last_name")  
    private String lastName;  
  
    private String street;  
  
    private String city;  
  
    private String state;  
  
    private String zip;  
  
    public String getOrderNumber() { return orderNumber; }  
  
    public OrderCustomerMessage setOrderNumber(String orderNumber) {...}  
  
    public String getCustomerEmail() { return customerEmail; }
```

# Sink Binding

```
spring.cloud.stream.bindings:  
  input:  
    destination: order_customers  
    group: order_to_customer_processor
```

```
public interface Sink {  
  
    String INPUT = "input";  
  
    @Input(Sink.INPUT)  
    SubscribableChannel input();  
  
}
```



# Sink and Source - a Processor

```
@EnableBinding(EmailProcessor.class)
public class OrderEmailProcessor {

    @Autowired
    JavaMailSender emailSender;

    @StreamListener(EmailProcessor.ORDER_CUSTOMERS_TO_EMAIL)
    @SendTo(EmailProcessor.EMAILS_SENT)
    public EmailSent sendOrderEmail(OrderCustomerMessage orderCustomerMessage){
        SimpleMailMessage message = new SimpleMailMessage();
        message.setTo(orderCustomerMessage.getCustomerEmail());
        message.setSubject("Thank You!");
        message.setText("Dear " + orderCustomerMessage.getFirstName() + ", thank you for ordering with us!");
        emailSender.send(message);

        EmailSent emailSent = new EmailSent();
        emailSent.setEmail(orderCustomerMessage.getCustomerEmail());
        emailSent.setEmailType(EmailSent.EmailType.ORDER_RECEIVED_EMAIL);
        return emailSent;
    }
}
```

# EmailProcessor Binding

```
public interface EmailProcessor {  
    String ORDER_CUSTOMERS_TO_EMAIL = "orderCustomersToEmail";  
    String EMAILS_SENT = "emailsSent";  
  
    @Input(ORDER_CUSTOMERS_TO_EMAIL)  
    SubscribableChannel orderCustomersToEmail();  
  
    @Output(EMAILS_SENT)  
    MessageChannel emailsSent();  
}
```

```
} spring.cloud.stream.bindings:  
  input:  
    destination: order_customers  
    group: order_to_customer_processor  
  orderCustomersToEmail:  
    destination: order_customers  
    group: email_processor  
  emailsSent:  
    destination: emails_sent
```

# Lastly a file Sink

```
java -jar file-sink-kafka-2.1.0.RC1.jar \  
--server.port=0 \  
--spring.cloud.stream.bindings.input.destination=order_customers \  
--file.mode=APPEND \  
--file.directory=output \  
--file.name=communications.log
```

File location and configuration

Destination binding for Kafka

# Demo Time





# Resources

[Kafka Tool](#)